

The Life and Adventures of LLVM

From Bytecode to the Executables

Rithik Sharma, PhD Student

Motivation?

Motivation?

- How does this talk align with the compiler class?

Motivation?

- How does this talk align with the compiler class?



Motivation?

- How does this talk align with the compiler class?




Motivation?

- How does this talk align with the compiler class?




Motivation?

- How does this talk align with the compiler class?



Motivation?

- How does this talk align with the compiler class?



Motivation?

- How does this talk align with the compiler class?



Motivation?

- How does this talk align with the compiler class?



Techniques
used by
compilers

Motivation?

- How does this talk align with the compiler class?

Motivation?

- How does this talk align with the compiler class?
- What are some shortcomings of early compilers?

Motivation?

- What are some shortcomings of early compilers?

Motivation?

- What are some shortcomings of early compilers?
 - Performance

Motivation?

- What are some shortcomings of early compilers?
 - Performance
 - Re-usability

Motivation?

- What are some shortcomings of early compilers?
 - Performance
 - Re-usability
 - Optimizations

Motivation?

- **What are some shortcomings of early compilers?**
 - **Performance**
 - **Re-usability**
 - **Optimizations**
 - **Correctness**

Motivation?

- **What are some shortcomings of early compilers?**
 - **Performance**
 - **Re-usability**
 - **Optimizations**
 - **Correctness**
 - **Scaling**

Motivation?

- **What are some shortcomings of early compilers?**



**Even Bob the Builder, is
confused about where to start**

Motivation?

- What are some shortcomings of early compilers?

We need a modern compiler!

Motivation?


- **Introduction to LLVM**

LLVM

LLVM

LLVM

Low Level Virtual Machine



2023
EURO LLVM
DEVELOPERS' MEETING













Motivation?

- **Introduction to LLVM**




Picture credits: Bob the Builder

Motivation?

- Introduction to LLVM

What happens inside
the front end?

- Lexical Analysis



LLVM

- **Lexical Analysis (tokenization or scanning)**
 - It breaks the source code into individual tokens, such as identifiers, keywords, literals, and operators.
 - Example of lexical analysis for a simple arithmetic expression:
"5 + 3 * (7 - 2)"

Token: Integer Value: 5

Token: Operator Value: +

Token: Integer Value: 3

Token: Operator Value: *

Token: Left Parenthesis

Value: (

Token: Integer Value: 7

Token: Operator Value: -

Token: Integer Value: 2


Token: Right Parenthesis Value:)

Motivation?

- Introduction to LLVM

**What happens inside
the front end?**

- Lexical Analysis
- Syntax Analysis



LLVM

- **Syntax Analysis**

- It builds the abstract syntax tree (AST) from the tokens.
- AST represents the hierarchical structure of the source code.
- Capturing the relationships between different elements and their corresponding expressions, statements, and declarations.

expr -> term
expr -> expr + term
term -> factor
term -> term * factor
factor -> Integer
factor -> (expr)


expr
/ \
term +
/ \ |
factor term
/ \
factor *
|
factor

Motivation?

- Introduction to LLVM

**What happens inside
the front end?**

- Lexical Analysis
- Syntax Analysis
- Semantic Analysis




LLVM

- **Semantic Analysis**
 - Semantic analysis ensures the program is well-formed and meaningful according to the language's rules and specifications.
 - It helps catch errors and inconsistencies that may not be detected during lexical and syntax analysis alone.

Motivation?

- **Introduction to LLVM**



**What happens inside
the front end?**

- **Lexical Analysis**
- **Syntax Analysis**
- **Semantic Analysis**
- **LLVM IR generation**

LLVM

- LLVM IR generation.

```
#include <iostream>

int main() {
    int x = 5;
    int y = 10;

    int z = x + y;

    return 0;
}
```

y

C++ Code

LLVM


- LLVM IR generation.

```
; Function Attrs: mustprogress noinline norecurse nounwind optnone uwtable
define dso_local noundef i32 @main() #4 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 5, i32* %2, align 4
    store i32 10, i32* %3, align 4
    %5 = load i32, i32* %2, align 4
    %6 = load i32, i32* %3, align 4
    %7 = add nsw i32 %5, %6
    store i32 %7, i32* %4, align 4
    ret i32 0
}
```

Motivation?

- **Introduction to LLVM**

**What happens inside
the front end?**



- **Lexical Analysis**
- **Syntax Analysis**
- **Semantic Analysis**
- **LLVM IR generation**
- **Optional
Optimizations**

LLVM

- **Optional Optimizations**
 - Constant folding - simplifies expressions involving constants and replaces them with their computed values.

```
; Function Attrs: mustprogress noinline norecurse nounwind optnone uwtable
define dso_local noundef i32 @main() #4 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 5, i32* %2, align 4
    store i32 10, i32* %3, align 4
    %5 = load i32, i32* %2, align 4
    %6 = load i32, i32* %3, align 4
    %7 = add nsw i32 %5, %6
    store i32 %7, i32* %4, align 4
    ret i32 0
}
```


LLVM

- **Optional Optimizations**
 - Constant folding - simplifies expressions involving constants and replaces them with their computed values.

```
; Function Attrs: mustprogress noinline norecurse nounwind optnone uwtable
define dso_local noundef i32 @main() #4 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %3 = alloca i32, align 4
    %4 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    store i32 5, i32* %2, align 4
    store i32 10, i32* %3, align 4
    %7 = add nsw i32 5, 10
    store i32 %7, i32* %4, align 4
    ret i32 0
}
```

Motivation?

- **Introduction to LLVM**



What happens inside the front end?

- **Lexical Analysis**
- **Syntax Analysis**
- **Semantic Analysis**
- **LLVM IR generation**
- **Optimizations**
- **Warnings and errors**

LLVM

- Warnings and errors.

```
#include <iostream>

int main() {
    int x = 5;
    int y = 10;

    int z = x * y; // Warning: Unused variable

    return 0; // Error: Missing semicolon
}
```

LLVM

- Warnings and errors.

```
program.cpp:7:9: warning: unused variable 'z' [-Wunused-variable]
    int z = x * y; // Warning: Unused variable
    ^
1 warning generated.
```


```
program.cpp:9:13: error: expected ';' after return statement
    return 0 // Error: Missing semicolon
    ^
;
1 error generated.
```

Motivation?

- **Introduction to LLVM**

What happens inside the middle end?

- **Data Flow Analysis (DFA)**



LLVM

- Data Flow Analysis (DFA)

```
#include <iostream>

int main() {
    int x = 10;
    int y = 5;
    int z;

    if (x > y) {
        z = x + y;
    } else {
        z = x - y;
    }

    return 0;
}
```

C++ code

LLVM


- Data Flow Analysis (DFA)

```
#include <iostream>

int main() {
    int x = 10;
    int y = 5;
    int z;

    if (x > y) {
        z = x + y;
    } else {
        z = x - y;
    }

    return 0;
}
```



C++ code

LLVM


- Is there a dead code?

```
#include <iostream>

int main() {
    int x = 10;
    int y = 5;
    int z;

    if (x > y) {
        z = x + y;
    } else {
        z = x - y;
    }

    return 0;
}
```



C++ code

LLVM


- Yes!

```
#include <iostream>

int main() {
    int x = 10;
    int y = 5;
    int z;

    if (x > y) {
        z = x + y;
    } else {
        z = x - y;
    }


    return 0;
}
```



Motivation?

- **Introduction to LLVM**

What happens inside the middle end?




- **Data Flow Analysis (DFA)**
- **Control Flow Analysis (CFA)**

Motivation?

- **Introduction to LLVM**

What happens inside the middle end?




- **Data Flow Analysis (DFA)**
- **Control Flow Analysis (CFA)**
- **Alias Analysis (AA)**

Motivation?

- **Introduction to LLVM**

What happens inside the middle end?


- **Data Flow Analysis (DFA)**
- **Control Flow Analysis (CFA)**
- **Alias Analysis (AA)**
- **Data Dependence Analysis (DDA)**



Motivation?

- **Introduction to LLVM**

What happens inside the middle end?




- **Optimizations**
 - Transformation passes
 - Analysis passes

Motivation?

- **Introduction to LLVM**

What happens inside the middle end?

- **Optimizations**
- **Optimization ordering**




LLVM

- Optimization ordering
 - The reason behind ordering?
 - What if there is a functionality change?
 - Transformation and Analysis passes.

Motivation?

- **Introduction to LLVM**


What happens inside the middle end?



- **Optimizations**
 - Optimization ordering
 - Generating optimized LLVM-IR

Motivation?

- **Introduction to LLVM**




Motivation?

- **Introduction to LLVM**


What happens inside the back end?

- **Target-Specific Code Generation**



Motivation?


- Target-Specific Code Generation



Motivation?

- **Introduction to LLVM**

What happens inside the back end?



- **Target-Specific Code Generation**
- **Instruction Selection**

Motivation?

- Instruction Selection
 - Instruction Matching
 - Cost analysis and pattern matching
 - DAG (Directed Acyclic Graph)

Motivation?


- Instruction Selection
 - **DAG (Directed Acyclic Graph)**
 - **DAG (Directed Acyclic Graph)** - It captures the dependencies and operations of the program as nodes and edges in a directed graph. Each node in the DAG represents an operation or value, and the edges represent the data flow between them.
 - Overall, DAG-based instruction selection in LLVM's backend is crucial in mapping high-level IR to target-specific machine code, enabling efficient and optimized code generation for a wide range of target architectures.

Motivation?

- Instruction Selection
 - **Global ISEL (Global Instruction Selection)**
 - It aims to improve instruction selection by performing the selection process across the entire function or module globally rather than on a per-basic-block basis, as done in the DAG approach.
 - Improved Code Quality: By considering a broader context and optimizing across the entire function or module.
 - Code Sharing: Global ISEL can identify opportunities for code sharing and reuse across different basic blocks and paths, leading to reduced code size and improved cache locality.
 - Simplified Code Generation: With Global ISEL, the instruction selection process becomes more unified and cohesive since it operates globally.

Motivation?

- **Introduction to LLVM**



What happens inside the back end?


- **Target-Specific Code Generation**
- **Instruction Selection**
- **Instruction Scheduling**

Motivation?

- **Instruction Scheduling**
 - The backend determines the order of instructions to be executed to maximize the performance of the generated code.
 - Instruction scheduling considers factors such as instruction dependencies, pipeline hazards, and the target architecture's specific execution characteristics to minimize stalls and improve instruction-level parallelism.

Motivation?

- **Introduction to LLVM**



What happens inside the back end?


- **Target-Specific Code Generation**
- **Instruction Selection**
- **Instruction Scheduling**
- **Register Allocation**

Motivation?

- **Register Allocation**
 - **Virtual Register Allocation** - virtual registers are initially unlimited and can hold any value.
 - This allows for efficient analysis and optimization without the limitations of physical registers.
 - **Register Interference Analysis** - determine which virtual registers may conflict with each other.

Motivation?

- **Introduction to LLVM**




What happens inside the back end?

- **Target-Specific Code Generation**
- **Instruction Selection**
- **Instruction Scheduling**
- **Register Allocation**
- **Code emission**


Motivation?

- **Introduction to LLVM**



Motivation?

- **Introduction to LLVM**



Motivation?

- Benefits of LLVM
 - Modularity and Extensibility

Motivation?

- Benefits of LLVM
 - Modularity and Extensibility
 - Portability

Motivation?

- Benefits of LLVM
 - Modularity and Extensibility
 - Portability
 - Optimizations

Motivation?

- Benefits of LLVM
 - Modularity and Extensibility
 - Portability
 - Optimizations
 - Just in time, execute the code on the fly

Motivation?

- Benefits of LLVM
 - Modularity and Extensibility
 - Portability
 - Optimizations
 - Just in time, execute the code on the fly
 - Supported tools (LLDB, GDB)

Motivation?

- Benefits of LLVM
 - Modularity and Extensibility
 - Portability
 - Optimizations
 - Just in time, execute the code on the fly
 - Supported tools (LLDB, GDB)
 - Community and easier Adoption

Thank
you :)